

Connected Vehicle Intelligence Hackathon

Solution Document Template

Submission Format: PDF (export of this document)

To be Submitted by: Jeevan Sarasram S S (solo entry)

Team Members & Roles: Jeevan Sarasram S S (Reg. No. RA2311056010035), solo developer: problem framing, architecture, backend, data engineering, ML, DevOps and testing; js9882@srmist.edu.in, jeevansiva2005@gmail.com

Problem Space Chosen: Predictive maintenance for mixed ICE / hybrid / EV fleets (with real-time critical-fault alerting)

Repository URL: <https://github.com/JeevanSarasramSS/fleetpulse> (tag v1.0-submission)

Demo Video URL: https://drive.google.com/file/d/1R0No_tNViiExbQ-IPfvOw9y1ivHt7P9Z/view?usp=sharing

Date of Submission: 01/10/2026

Table of Contents

- [1. Executive Summary](#)
- [2. Problem Statement & Validation](#)
 - [2.1 Problem Statement](#)
 - [2.2 Evidence & Validation](#)
 - [2.3 Impact & Success Metrics](#)
- [3. Solution Description](#)
 - [3.1 Solution Overview & User Journey](#)
 - [3.2 Key Value Proposition](#)
 - [3.3 Innovative Ideas](#)
- [4. Feature List](#)
- [5. Solution Architecture \(High-Level Design\)](#)
 - [5.1 Architecture Overview](#)
 - [5.2 Technology Stack & Justification](#)
 - [5.3 Data Architecture](#)
 - [5.4 Deployment View](#)
- [6. Low-Level Design](#)
 - [6.1 Layering & Separation of Concerns](#)
 - [6.2 Design Principles Applied](#)
 - [6.3 Design Patterns Used](#)
 - [6.4 Interfaces, Contracts & Runtime Flows](#)
 - [6.5 Algorithms & Data Structures](#)
- [7. Non-Functional Requirements & Performance Benchmarks](#)
- [8. Security & Compliance](#)
- [9. Test Strategy](#)
- [10. Observability](#)
- [11. AI / ML Component \(if used\)](#)
- [12. Architecture Decisions, Risks & Future Enhancements](#)
- [13. Demo Video \(5 Minutes Maximum\)](#)
- [14. Repository Checklist](#)
- [15. Conclusion](#)
- [16. Declarations](#)
- [17. Appendix \(if any\)](#)

1. Executive Summary

Brief overview (half a page maximum) of:

- The problem you chose and who it affects
- What your solution does, in one or two sentences
- Key results achieved (throughput, latency, impact numbers)
- Unique features, innovations or optimisations you implemented

Problem. Unplanned breakdowns are the most expensive event in a commercial fleet: a tow, an emergency repair and two days of a vehicle earning nothing. Fleet managers running mixed petrol, hybrid and EV fleets see fault codes and warning lights only after the fact, scattered across OEM portals with different formats.

Solution. FleetPulse streams telemetry from every vehicle, raises critical faults on the manager's screen within seconds, and every minute ranks the whole fleet by its probability of breaking down in the next 7 days, with the reasons and a dollar estimate. A fleet copilot answers questions over the same data and proposes work orders that a human approves.

Results (measured on the submitted code). 100,000 simulated vehicles across 3 tenants and 2 OEM payload formats; load-tested end to end at up to ~25,000 events/s with 3x bursts and zero data loss (every unique event written), ~5,000 events/s in the default demo (CPU hot path 33.9K events/s per core); data freshness vehicle → dashboard ~0.15 s; critical alert vehicle → screen p50 0.17 s and p95 2.6 s even through a 3x burst (target < 5 s); 100,000 vehicles scored in ~10 s; risk model ROC-AUC 0.878 vs 0.688 for the mileage-since-service baseline and 4.2x the baseline's precision on the top 2% of the fleet; API p95 110 ms at 50 concurrent users; slowest query 174 ms → 0.58 ms; recovery from killing the processors and restarting the broker in ~26 s with zero data loss; the production Kubernetes manifests deploy unchanged (config only) on a kind cluster in CI.

What is distinctive. Effectively-once processing without Kafka transactions (at-least-once + idempotent sinks); explainable per-vehicle risk factors; a copilot with propose-only tools, tenant-scoped data access, injection screening, VIN grounding checks and an append-only audit trail; and honest gap reporting (section 12).

2. Problem Statement & Validation

2.1 Problem Statement

- **Format:** “[User] needs a way to [job to be done] because [pain], which today costs [measurable impact].”
- **Primary User:** Fleet manager, dealer, lender, driver or OEM — name one
- **Secondary Stakeholders:** Others affected by the problem

“A fleet manager needs a way to know which vehicles will fail in the next week and act before they do, because breakdowns are discovered only when a driver is stranded, which today costs roughly

US\$2,400 per incident in towing, emergency repair and two days of lost vehicle utilisation (our working assumption, see 2.2).”

Primary user: Fleet manager of a commercial or rental fleet (1,000–100,000 vehicles, mixed ICE / hybrid / EV).

Secondary stakeholders: drivers (safety, fewer roadside failures), maintenance workshops (planned rather than emergency work), fleet finance (cost and utilisation), OEMs (warranty and fault-pattern insight), lenders and insurers (asset condition).

2.2 Evidence & Validation

- **Facts vs Assumptions:** Separate what you know from what you assume
- **Validation Method:** Published study, public dataset, simulation, or interviews
- **Existing Alternatives:** Aftermarket dongles, OEM portals, Motorq Fuse, other telematics tools — and why they fall short

Evidence / Assumption	Source or Method	What It Shows	Confidence
Unplanned breakdowns cost far more than planned repairs (tow, emergency labour, downtime)	Industry fleet-maintenance literature; our assumption US\$2,400 vs US\$650	Size of the pain per incident	Medium
Connected cars will generate up to ~25 GB/hour; ~95% of new cars connected by 2030	Problem statement (McKinsey / S&P)	Data volume requires stream processing, not batch exports	High
Rising coolant temperature and falling 12 V voltage precede failures by days	Simulation (hidden-wear model) and common maintenance practice	Symptoms carry a learnable early signal	Medium
Mileage-since-service rule is the common baseline today	Assumption from OEM service schedules	Baseline the model must beat	Medium
Our model ranks risk 4.2x better than the rule at the top 2%	Back-test on 25,000 held-out simulated vehicles	Where to send scarce workshop capacity	High (on simulated data)

Facts are taken from the brief and public sources; numbers we could not verify are marked as assumptions and are the inputs to our savings estimate, so they can be changed in one place (fleetpulse/agent/copilot.py).

Validation method: a simulation in which every vehicle has a hidden wear state that drives observable symptoms (coolant temperature, 12 V voltage, fault-code rate) and a breakdown hazard. The model is trained and back-tested on 100,000 simulated vehicle-weeks, and the same simulator feeds the live system.

Existing alternatives: aftermarket OBD dongles (hardware cost and install, one data format); OEM portals (one brand per portal, alerts after the fault); rule-only telematics (alerts on a code, no ranking of who fails next); Motorq Fuse (AI recommendations on OEM data, commercial). FleetPulse combines device-free multi-OEM ingestion, sub-5-second alerts and an explainable 7-day forecast in one view.

2.3 Impact & Success Metrics

Metric	Baseline Today	Target	How Measured / Estimated
Critical fault to manager's screen	Hours to days (driver call)	< 5 s	Measured p50 0.17 s, p95 2.6 s via WebSocket timestamps
Precision of 'service now' list (top 2%)	6.6% (mileage rule)	> 20%	Back-test: 27.8%
Unplanned breakdowns / 1,000 vehicles / week	28 (simulated base rate)	-30%	Estimated: servicing top-50 avoids ~31 of the tenant's expected breakdowns
Net weekly saving, one 33K-vehicle tenant	US\$0	> US\$50K	Copilot estimate: US\$63K net

- **Scale of Impact:** Who benefits, by how much, at 10K vs 100K vehicles
- **Wider Impact:** Safety, cost, environment or compliance, where relevant

Scale of impact: With a 2.8% weekly breakdown base rate in the simulation, a 10K-vehicle fleet sees ~280 breakdowns a week and a 100K fleet ~2,800. Servicing the 50 highest-risk vehicles of one 33K-vehicle tenant avoids an expected 31 breakdowns this week: US\$74.5K avoided against US\$11.4K planned spend, net US\$63K per week (live copilot estimate from the running system).

Wider impact: safety (critical faults such as HV-isolation and coolant over-temperature are surfaced in seconds), cost (planned repairs instead of tows), environment (misfire and catalyst faults fixed earlier cut emissions), compliance (audit trail, masked location for analysts, erasure flow).



3. Solution Description

3.1 Solution Overview & User Journey

- **What It Does:** Plain-language description of the solution
- **User Journey:** Step by step, e.g. vehicle event → detection → alert → action → outcome
- **Screenshot:** At least one screenshot of the working product

FleetPulse is a web dashboard backed by a streaming pipeline. Vehicles (here, the simulator) stream telemetry; FleetPulse normalises every OEM format into one event, detects faults in real time, scores every vehicle's 7-day breakdown risk each minute, and lets a fleet manager act from one screen.

User journey: (1) a Borealis van reports coolant at 119 °C and code P0217 in OEM-specific JSON; (2) the processor normalises it, validates the VIN check digit, drops duplicates and evaluates sliding-window rules; (3) an OVERHEAT_TREND / CRITICAL_DTC alert is written idempotently and pushed over WebSocket, appearing on the dashboard ~0.2 s after the vehicle timestamp; (4) the manager opens the vehicle: 73% 7-day risk, reasons “overdue for service, coolant running hot, 12 V battery voltage low”, the coolant trend and recent alerts; (5) they click Propose work order (or ask the copilot); (6) the work order appears as proposed and a manager approves it; every step is in the audit log.



3.2 Key Value Proposition

- **Customer Job:** What the user is trying to get done

- **Pain Relieved:** What gets easier, cheaper or safer
- **Gain Created:** New capability the user did not have
- **Differentiation:** Why your solution, not the alternatives

Customer job: keep every vehicle earning by fixing problems before they strand a driver.

Pain relieved: no more discovering failures from the roadside; one place for all OEMs; no manual triage of thousands of fault codes.

Gain created: a ranked, explained list of who will fail this week, a dollar figure for acting now, and a copilot that turns a question into a proposed work order.

Differentiation: device-free and multi-OEM by design (adapter per OEM), explainable ML that beats the mileage rule fleets use today by 4.2x on the top of the list, and AI that can only propose, never act, with a full audit trail.

3.3 Innovative Ideas

- **Up to Three Ideas:** A technique, data insight, user experience or business model
- **What Is New:** Compared against prior work you studied
- **Evidence:** Proof that each idea works

1. Effectively-once without distributed transactions. At-least-once Kafka consumption with offsets committed only after the database transaction, plus idempotent sinks (unique (vin, seq, ts) on telemetry, a deterministic dedup_key per vehicle/rule/code/hour on alerts) and a Bloom filter fast path. Evidence: 1% duplicate and 2% out-of-order events are injected continuously; killing both processors and restarting the broker produced no duplicate alerts and the backlog drained at ~9K events/s.

2. Symptom-to-hazard simulation as a validation harness. The simulator gives each vehicle a hidden wear level that drives both noisy symptoms and the breakdown hazard, so the model must infer wear from symptoms, as in reality. This lets us measure model lift against the industry rule (service overdue) honestly: ROC-AUC 0.878 vs 0.688.

3. Propose-only, grounded copilot. Tools are tenant-scoped SQL, the only write tool creates a 'proposed' work order, every VIN in an answer must come from a tool result (otherwise the answer is withheld), injection attempts are refused, and every tool call is audited. It runs deterministically offline and uses Claude tool-use when a key is configured.

4. Feature List

Every feature must be traceable to code. Priority uses MoSCoW (Must / Should / Could / Won't). Status: Done / Partial / Planned. Give the demo video timestamp for Done features.

ID	Feature	User Story	Priority	Status	Code Path	Video
F-01	Multi-OEM telemetry ingestion and normalisation	As a platform team I want new OEM formats handled by an adapter so that onboarding needs no downtime	Must	Done	fleetpulse/core/normalise.py	[mm:ss]
F-02	100K-vehicle simulator with bursts, duplicates, late events	As an engineer I want realistic load so that we can prove scale	Must	Done	fleetpulse/simulator/	[mm:ss]
F-03	Real-time critical fault alerts (< 5 s)	As a fleet manager I want critical faults instantly so that drivers are not stranded	Must	Done	fleetpulse/processor/	[mm:ss]
F-04	7-day breakdown risk ranking with reasons	As a fleet manager I want to know who will fail this week and why so that I book them in	Must	Done	fleetpulse/ml/	[mm:ss]
F-05	Live fleet map	As a dispatcher I want to see where vehicles are so that I can route recovery	Should	Done	web/index.html, /api/v1/live	[mm:ss]

ID	Feature	User Story	Priority	Status	Code Path	Video
F-06	Fleet copilot with work-order proposals	As a fleet manager I want to ask questions in plain English and get actions I approve	Should	Done	fleetpulse/agent/	[mm:ss]
F-07	RBAC, tenant isolation, location masking	As an admin I want analysts to see only coarse locations of our own fleet	Must	Done	fleetpulse/api/, core/security.py	[mm:ss]
F-08	Audit log and right-to-erasure	As a DPO I want every access logged and driver PII erasable	Must	Done	db/001_schema.sql, DELETE /drivers	[mm:ss]
F-09	Savings estimate	As finance I want the dollar impact of acting now	Could	Done	copilot.estimate_savings	[mm:ss]
F-10	OIDC SSO and device mTLS	As security I want federated login and device identity	Should	Planned	docs/threat-model.md	—
F-11	ClickHouse/Parquet cold tier	As a platform team I want cheap years of history	Could	Planned	ADR-0002	—

5. Solution Architecture (High-Level Design)


5.1 Architecture Overview

- **System Context Diagram:** C4 level 1 — your system, users and external systems (OEM clouds, identity provider, LLM API, data warehouse)
- **Container Diagram:** C4 level 2 — every service, data store and broker, with the protocol on each arrow (MQTT, HTTPS, gRPC, Kafka)
- **Data Flow:** Path of one telemetry event from vehicle to insight, with expected latency per hop

System context (C4 L1): users (fleet manager, analyst, admin) use the FleetPulse web app; FleetPulse ingests from OEM clouds/vehicles (simulated), optionally calls an LLM API (Anthropic) for copilot planning, and would authenticate through an OIDC provider in production.

Container view (C4 L2) and data flow: see diagram. One telemetry event: vehicle → Kafka (~20 ms producer linger) → processor batch (0.1–0.2 s; alerts are written and pushed before the bulk telemetry upsert) → Redis pub/sub → WebSocket (< 50 ms) → screen. Measured p50 0.17 s, p95 2.6 s vehicle timestamp to screen across a 3x burst.

FleetPulse: container view (C4 level 2) with per-hop latency



Vehicle → screen critical-alert latency measured p50 1.38 s (target < 5 s). Same images run on docker compose, kind/k3d, EKS, GKE or AKS: endpoints come from env/Secrets only.

5.2 Technology Stack & Justification


Layer	Choice	Why This, and What You Rejected
Ingestion / Messaging	Kafka protocol (Redpanda locally, MSK/Confluent in cloud), key = VIN	Replay, partitioned per-vehicle ordering, consumer groups; RabbitMQ rejected (no replay, lower throughput)
Stream / Batch Processing	Python consumer with batch upserts; batch scorer + incremental rollup	Simple, testable and measured at 34K ev/s/core; Flink/Spark rejected as too heavy for the time box (next step at scale)
Relational / NoSQL / Cache / Search / Vector	PostgreSQL 16 (3NF + partitions) · Redis 7 · pgvector HNSW	ACID core and analytics in one engine; Redis for sub-ms live state/geo; pgvector avoids another system. ClickHouse/Scylla deferred (ADR-0002)
Backend / Frontend	FastAPI (async, OpenAPI) · single-page HTML/JS + Leaflet	Fast to build, typed contracts; no front-end build step; React rejected for time
ML / AI	scikit-learn HistGradientBoosting · Claude tool-use (optional)	Strong tabular baseline with explanations; LLM only plans tool calls, never owns facts
Infrastructure / CI-CD / Observability	Docker compose, Kubernetes manifests, Terraform (AWS), GitHub Actions, Prometheus + Grafana	Portable across clouds by configuration (proven on kind in CI); CI runs lint, unit tests with coverage gate, SAST, e2e (integration, contract, BDD), ZAP DAST, image scan and chaos

5.3 Data Architecture

- **ER Diagram:** Relational core normalised to 3NF; list any deliberate denormalisation and why
- **Polyglot Map:** Which data lives in which store, with its CAP choice (CP or AP)
- **Capacity Estimate:** Events/sec, bytes/event, daily and yearly volume, partition keys, hot / warm / cold retention
- **Query Optimisation:** Top three queries with EXPLAIN ANALYZE before and after

Query	Before (ms)	After (ms)	Change Made
Vehicle detail: last 30 min of telemetry for one VIN (3M rows)	174	0.58	Composite index (vin, ts DESC) per daily partition + partition pruning on ts
Top-25 risk for a tenant (100K scores)	27.3	0.40	Index (tenant_id, score DESC) → index scan with LIMIT, no sort
Open alerts feed, keyset page	2.7	0.15	Partial index WHERE acked_at IS NULL + keyset on alert_id instead of OFFSET

Relational core (3NF) · arrows point from child (FK) to parent



Telemetry (partitioned by hour, unique (vin, seq, ts)), the vehicle_daily incremental rollup and fault_knowledge (pgvector) sit outside the 3NF core. Deliberate denormalisation: alert.tenant_id and risk_score.tenant_id avoid a 3-table join on the hottest tenant-filtered queries.

Polyglot map: Postgres 3NF core: tenants, users, subscriptions, fleets, vehicles, drivers, alerts, work orders, risk scores, audit (CP). Postgres partitioned telemetry and vehicle_daily rollup (writes buffered by Kafka, AP end to end). Redis: live state, GEO index, rate limits, pub/sub, cache (AP). pgvector: fault knowledge embeddings (CP).

Capacity: 100K vehicles × 1 event/s = 100K events/s, ~0.5 KB JSON (~120 B per stored row) → 8.6 TB/day raw JSON, ~1 TB/day as rows; ~0.3 TB/day compressed in a columnar cold tier. Kafka partition key VIN (uniform hash, no hot spots); telemetry partitioned by hour with (vin, ts) indexes. Hot: Redis latest state (minutes) and raw Postgres partitions (2 h locally, configurable). Warm: Postgres/ClickHouse 90 days (~25 TB compressed at full scale). Cold: Parquet on S3 7 years (~US\$0.023/GB-month → ~US\$2.5K/month per PB-year).

Denormalisation: alert.tenant_id and risk_score.tenant_id duplicate fleet.tenant_id to serve the hottest tenant-filtered queries from one index; consistency is guaranteed because both are written from the same vehicle → fleet lookup.

5.4 Deployment View

- **Deployment Diagram:** Clusters, namespaces, replicas, autoscaling, networking, secrets
- **Cloud-Agnostic Approach:** How the same deployment runs on a second cloud or locally

Kubernetes namespace fleetpulse (pod-security restricted): api Deployment (3 replicas, HPA 3–20 on CPU, PDB minAvailable 2, zone topology spread), processor Deployment (6 replicas ≤ partitions; KEDA

on consumer lag in production), risk-scoring CronJob every 5 minutes, default-deny NetworkPolicy, TLS Ingress via cert-manager. Containers run as non-root with read-only root filesystems and dropped capabilities. Secrets come from a K8s Secret synced from Vault / cloud secret manager (infra/k8s/fleetpulse.yaml).

Portability proof: a CI job deploys these exact manifests on a kind Kubernetes cluster through a kustomize overlay that changes configuration only (image, Secret, replica counts; Ingress omitted), under the restricted Pod Security profile, then smoke-tests login and live streaming through the Service. Cloud-agnostic approach: the application only reads KAFKA_BOOTSTRAP, PG_DSN, REDIS_URL and JWT_SECRET from the environment (12-factor). Terraform for AWS (infra/terraform/aws: VPC, EKS, MSK, RDS Multi-AZ, ElastiCache, all encrypted) provisions the backing services; the same manifests run on GKE/AKS with Confluent/Event Hubs, Cloud SQL/Flexible Server and Memorystore/Azure Cache, or locally on kind/k3d and docker compose, with no code change.

6. Low-Level Design

6.1 Layering & Separation of Concerns

- **Architecture Style:** Layered, hexagonal (ports and adapters) or clean architecture — justify
- **Layer Diagram:** Internal layers of at least one core service and the allowed dependency direction
- **Folder Structure:** Repository tree, two levels deep

Layer	Responsibility	Must Not
Presentation / API	fleetpulse/api: HTTP, WebSocket, validation (Pydantic), auth, RBAC, DTO mapping	Contain business rules
Application / Service	processor/run.py, ml/score.py, agent/copilot.py: use cases, transactions, orchestration	Depend on a specific OEM format
Domain	fleetpulse/core + processor/rules.py: canonical event, VIN/DTC rules, algorithms, alert rules (pure Python, no I/O)	Import framework or infrastructure code
Infrastructure	api/deps.py, Kafka/Postgres/Redis clients, SQL	Leak vendor types into the domain

Style: layered with ports-and-adapters at the edges. The domain (core, rules) is pure Python with no I/O, so it is unit-tested to 99% coverage and reused by the processor, the batch scorer and the API. Dependencies point inward: api → application → domain; infrastructure implements adapters (OEM adapters, Kafka, Postgres, Redis).

Folder structure: fleetpulse/{core, simulator, processor, ml, agent, api} · web/ · db/ · infra/{k8s, terraform/aws, grafana} · tests/{unit, integration, contract, acceptance, load} · docs/{adr, diagrams, evidence, solution} · .github/workflows/

6.2 Design Principles Applied

- SOLID: single responsibility per module (normalise, rules, streaming structures); open/closed for OEMs (add an adapter, touch nothing else); dependency inversion in the copilot (planner strategy over the same tools).
- 12-Factor App: all config from environment (fleetpulse/config.py); stateless API and processors (state in Kafka offsets, Postgres, Redis); disposable containers with SIGTERM handling; one image, many process types.
- Idempotency, Fail-Fast, Least Privilege, DRY, KISS: idempotent sinks and alert dedup keys; schema and VIN validation fail fast into the DLQ; non-root read-only containers and role-scoped endpoints; shared feature code for training and scoring (DRY); one Python image instead of five stacks (KISS).

6.3 Design Patterns Used

Pattern	Problem It Solves in Your System	Location in Code
Adapter	Maps each OEM's payload format to one canonical event	fleetpulse/core/normalise.py (AuroraAdapter, BorealisAdapter)
Repository / ports	Keeps SQL out of domain logic; rules and algorithms are pure and unit-tested	fleetpulse/processor/rules.py (pure) vs processor/run.py (I/O)
Idempotent consumer	At-least-once delivery without duplicate rows or alerts	processor/run.py, unique keys in db/001_schema.sql
CQRS (lite)	Writes go through Kafka/processor; reads from Redis hot state, the incremental vehicle_daily rollup and the risk table	processor/, ml/score.py, api/main.py
Strategy	Copilot planner: LLM tool-use or deterministic router behind one interface	fleetpulse/agent/copilot.py
Observer / pub-sub	Alerts fan out to WebSocket clients per tenant	Redis PUBLISH in processor, /ws/alerts in API
Graceful degradation	Rate limiter fails open; map falls back from Esri to OpenStreetMap tiles and works without any tile CDN; pollers survive API restarts; copilot falls back from LLM	api/deps.py, web/index.html, agent/copilot.py

6.4 Interfaces, Contracts & Runtime Flows

- **API Contract:** OpenAPI / AsyncAPI link; versioning, pagination, error format, rate limits

- **Event Schemas:** Topics, keys, schema format (Avro / Protobuf / JSON Schema), schema evolution
- **Sequence Diagrams:** Two critical flows, including one failure path (broker down, duplicate event, OEM format change)

API contract: OpenAPI 3 generated by FastAPI at /openapi.json (Swagger UI at /docs). Versioned under /api/v1. Keyset (cursor) pagination with opaque base64 cursors on vehicles and alerts. Errors are RFC 7807 application/problem+json. Rate limit 600 requests/min per user, HTTP 429 with Retry-After. WebSocket /ws/alerts?token=... streams tenant-scoped alerts.

Event schemas: topic telemetry.raw (12 partitions, key = VIN, retention 24 h); telemetry.dlq (reason header); alerts. Payloads are JSON in two OEM shapes (Aurora: flat ISO-time metric; Borealis: nested epoch imperial with a raw fault string), mapped to one CanonicalEvent dataclass. Evolution: a new OEM or version is a new adapter class registered in ADAPTERS; unknown shapes go to the DLQ rather than blocking a partition. Avro/Protobuf with a schema registry is the production next step.

Flow 1 (happy path, critical fault): simulator.produce(key=VIN) → processor.consume(batch≤5000) → normalise → is_valid_vin → bloom.add(vin:seq) → ReorderBuffer.observe → RuleEngine.evaluate → INSERT telemetry ... ON CONFLICT DO NOTHING; INSERT alert ... ON CONFLICT (dedup_key) DO NOTHING RETURNING → Redis HSET/GEOADD/PUBLISH → commit offsets → API WebSocket → dashboard.

Flow 2 (failure: duplicate + broker down): the device retries an event (duplicate) → Bloom filter drops it; if the Bloom state was lost in a crash, the unique index and dedup_key absorb the replay. If the broker is down, the producer buffers and applies back-pressure (BufferError → poll and retry); processors block on consume and resume from the last committed offset. Measured: processors killed and broker restarted at 13:02:10, consumption resumed and drained the backlog at ~9K events/s against ~5.9K/s input.

6.5 Algorithms & Data Structures

- **Problem Solved:** e.g. nearest reachable charger, trip segmentation, VIN / DTC parsing
- **Algorithm Used:** Graph, dynamic programming, regex, streaming structures — justify the choice
- **Pseudocode:** For major functions
- **Time & Space Complexity:** Annotate all
- **Scale Tested:** Input size and measured runtime

Problem	Algorithm / structure	Complexity	Scale tested
VIN validation	Regex <code>^[A-HJ-NPR-Z0-9]{17}\$</code> + ISO 3779 weighted check digit (mod 11)	$O(17)$ time, $O(1)$ space	100K generated VINs; 500 random round-trips in tests
DTC parsing from raw OEM strings	Regex <code>\b([PCBU])([0-3])([0-9A-F])([0-9A-F]{2})\b</code> with de-duplication	$O(n)$ in payload length	every event (~6K/s)
Duplicate detection	Bloom filter, k hashes by double hashing (blake2b), $m = -n \ln p / (\ln 2)^2$	$O(k)$ per op, ~9 MB for 5M keys at $p=0.001$	5M capacity per processor
Top-K fault codes	Count-Min Sketch ($w=2048, d=4$) + bounded candidate set	$O(d + K \log K)$ per update	live stream, shown in /stats
Overheat and harsh-brake patterns	Time-based sliding window with monotonic deque for max	amortised $O(1)$ per event	100K vehicles, bounded state
Late events	Per-VIN last-sequence map (never overwrite newer live state)	$O(1)$	2% out-of-order injected
Trip segmentation of noisy speed	2-state DP (Viterbi-style) with switch penalty	$O(n)$ time and space	unit-tested; blips do not split trips
Nearest reachable depot / charger	Dijkstra with binary heap and range cut-off	$O((V+E) \log V)$	unit-tested
Live map	Geohash (base32) and Redis GEO index (geohash-ordered sorted set)	$O(\log n + m)$	4,000 vehicles per map refresh
Pagination	Keyset on primary key	$O(\log n)$ per page at any depth	100K vehicles

Pseudocode (processor batch): for m in batch: $e \leftarrow \text{adapter}(m)$; if $\text{!valid}(e.\text{vin}) \rightarrow \text{DLQ}$; if $\text{bloom.add}(e.\text{vin}:e.\text{seq}) \rightarrow \text{skip}$; $\text{inOrder} \leftarrow \text{order.observe}(e)$; $\text{rows} += e$; $\text{alerts} += \text{rules}(e)$; if inOrder : $\text{latest}[e.\text{vin}] \leftarrow e$. Then one transaction: upsert rows , $\text{insert alerts RETURNING new}$; then Redis pipeline; then commit offsets.

Measured: hot path (parse → normalise → VIN → Bloom → rules) 200,000 events in 5.90 s = 33.9K events/s on one core (tests/load/processor_bench.py).

7. Non-Functional Requirements & Performance Benchmarks

NFR	Target (Case Study)	Achieved	How Measured
Ingest Throughput	100K+ events/sec; 3x burst for 5 min without loss	Zero loss at 12.8K/s and 25.7K/s offered with 3x bursts (rows written = unique events sent); Kafka accepted 69K/s, processing ceiling ~24K/s on one laptop Postgres (scale-out path in ADR-0002); 33.9K/s per core CPU path	tests/load/pipeline_load.py; docs/evidence/processor_bench.txt
End-to-End Latency	< 2 s dashboard; < 5 s critical alert	Data freshness vehicle → dashboard p50 ~0.15 s (KPIs poll every 1 s); critical alert p50 0.17 s, p95 2.6 s, max 3.2 s through a 3x burst	WebSocket push timestamp minus vehicle timestamp

NFR	Target (Case Study)	Achieved	How Measured
API Latency	p95 < 200 ms; p99 < 500 ms	p95 18 ms @10, p95 110 ms / p99 132 ms @50, p95 224 ms @100 concurrent users, with the pipeline running on the same box	tests/load/api_load.py
Resilience	Recovers after broker / pod failure	Processors killed + broker restarted: flow resumes ~26 s after the restart (10 s consumer session timeout), zero loss, no duplicates; proven in CI on every push	docs/evidence/chaos.txt
Availability	99.9%, no single point of failure	Design: 3 API replicas + PDB, RF=3 brokers, Multi-AZ RDS, Redis failover; not measured	infra/k8s, infra/terraform

- **Load Test Setup:** Tool, data volume, hardware / cluster size
- **Results:** Graphs or screenshots of throughput, p95 / p99 and consumer lag

Original load test setup: everything on one 4-vCPU cloud VM under docker compose (Redpanda, Postgres 16, Redis, 2 processors, simulator, batch, API) with 100K vehicles at a 5,000 events/s base rate and 3x bursts for 10 s every 2 minutes. API load: tests/load/api_load.py (async httpx) against 4 uvicorn workers, rate limiter raised for the test. Processor CPU benchmark: tests/load/processor_bench.py.

After hardening, everything was re-measured on a 16-vCPU laptop: alert latency with 4 processor replicas (docs/evidence/alert_latency.txt); end-to-end throughput with tests/load/pipeline_load.py (parallel producers, 3x bursts, zero-loss check); API latency with a multi-process client inside the compose network (api_load.txt); and a 45-minute soak (soak.txt). Raw outputs in docs/evidence/.

Results: API 2,000 requests at 10 concurrent: 326 rps, p50 25 ms, p95 71 ms, p99 110 ms (all 200). At 50 concurrent on the same shared box: p95 456 ms, p99 576 ms: the CPU is shared with the pipeline, and the fix is horizontal (API HPA, read replicas). Processor ~2.5K events/s per replica including Postgres writes; consumer lag drained after chaos at ~9K events/s.

8. Security & Compliance

- **Threat Model:** STRIDE for the ingestion path and public API; top five threats and controls
- **Authentication & Authorisation:** OAuth2 / OIDC, RBAC, tenant isolation
- **Device Identity & Encryption:** mTLS, TLS in transit, encryption at rest, secrets management
- **Privacy:** How location and driver data is minimised, masked, retained and erased (GDPR, DPDP Act 2023)
- **AI Safety (if applicable):** Prompt-injection defences, tool permissions, human approval, audit log

#	STRIDE threat	Control	Status
1	Spoofing a vehicle	VIN check digit + unknown-VIN DLQ; per-device mTLS at the edge	Done / mTLS planned
2	Tampering / replay flooding alerts	Bloom + unique (vin, seq, ts); alert dedup_key; per-VIN sequence	Done
3	Repudiation of approvals / access	Append-only audit_log (DB trigger) on reads, acks, approvals, erasure, logins, agent tools	Done
4	Cross-tenant disclosure	Tenant claim in JWT applied to every query; 404 for foreign VINs; tenant pub/sub channels	Done
5	DoS via bursts / API abuse	Kafka buffering, producer back-pressure, Redis rate limit (429), bounded rule state, HPA	Done
6	Privilege escalation via prompt injection	Injection screen, scoped tools, propose-only writes, human approval, grounding	Done

AuthN/AuthZ: JWT (HS256, 2 h expiry, issuer check) issued after PBKDF2-SHA256 (200K iterations) password verification; RBAC roles admin / fleet_manager / analyst enforced per endpoint; tenant isolation on every query. Production design: OIDC (Keycloak or cloud IdP) issuing RS256 tokens, same claims.

Device identity & encryption: designed for per-device X.509 mTLS at the MQTT/Kafka edge; TLS 1.3 at ingress (cert-manager); RDS, ElastiCache and MSK encrypted at rest (AES-256 KMS) and in transit (Terraform); secrets via K8s Secret synced from Vault; .env.example only in git. Security headers (nosniff, frame deny, no-referrer, no-store) on every response.

Privacy (GDPR, DPDP Act 2023): analysts see only geohash-5 (~5 km) masked positions; driver PII lives in one table and is erased by DELETE /api/v1/drivers/{id} (PII nulled, operational rows kept, audit row written); telemetry retention is a partition drop (24 h in Kafka, 7–90 days warm, then cold); Mumbai region for Indian data residency.

AI safety: see section 11.

9. Test Strategy


Test Type	Tools	No. of Tests	Coverage / Result	In CI?
Unit	pytest, pytest-cov	50	98% line coverage on all domain modules (algorithms, rules, guardrails, features, fleet generator); I/O endpoints covered by the live-stack suites	Yes (80% gate)
Integration & Contract	pytest + httpx against docker compose; consumer-driven contracts (dashboard as consumer) verified against the live API and its OpenAPI	8 + 10	auth, tenant isolation, RBAC, masking, pagination, erasure, guardrail, response contracts: all pass	Yes (e2e job)
Acceptance (BDD)	pytest-bdd Gherkin features (tests/acceptance/features)	6 scenarios	risk list, alert ack + audit, copilot work order + human approval, analyst masking, tenant isolation, erasure: all pass	Yes (e2e job)
Performance / Load / Soak	pipeline_load.py (zero-loss check, parallel producers, 3x bursts), api_load.py (multi-process client), soak.py (45 min), processor_bench.py	4	Zero loss at 25.7K ev/s offered; API p95 110 ms @50; soak 15.8M events, freshness p50 0.19 s, burst backlog grows after ~25 min (p95 7.9 s) as the hourly partition outgrows cache	Manual (needs a 16-vCPU host)
Security (SAST, DAST, Dependency, Image)	Semgrep, Bandit, pip-audit, Trivy, OWASP ZAP baseline	5	Bandit clean (0 medium/high); ZAP and Trivy reports in the CI run	Yes
Compliance & Chaos	erasure + audit test; kill processors + restart broker	2	Audit row on erase; flow resumes ~26 s after a broker restart, no data loss	Yes

- **Edge Cases Covered:** e.g. duplicate events, out-of-order data, unknown OEM format, broker outage
- **Sample Input / Output:** If helpful
- **Report Links:** Coverage, load-test and security scan reports
- **Edge cases covered:** duplicate events, out-of-order events, unknown OEM format (DLQ), invalid VIN check digit, lowercase VINs, unknown fault codes, EVs without coolant, empty inputs, expired/tampered JWT, cross-tenant VIN access, analyst masking, prompt injection, hallucinated VINs, unbounded state under VIN floods, broker and processor failure.
- **Sample:** Borealis payload {vehicleId, epoch, signals:{speedMph:10, coolantTempF:212, lvBatteryVolts:11.5}, faultString:'DTC:P0217;DTC:C0035'} → canonical {speed_kmh:16.09, coolant_c:100.0, dtcs:[P0217, C0035]} → alerts CRITICAL_DTC P0217 (sev 5), DTC C0035, LOW_12V.
- **Reports:** coverage (pytest --cov, uploaded by CI), docs/evidence/api_load.txt, processor_bench.txt, explain_analyze.txt, chaos.txt; Semgrep/Bandit/pip-audit/Trivy/ZAP output and the e2e JUnit report in the GitHub Actions run.

10. Observability

- **Dashboard Screenshot:** Throughput, consumer lag, error rate, latency percentiles
- **Troubleshooting Walk-Through:** How you would find the cause of a latency spike using logs, metrics and traces

Every service exposes Prometheus metrics: fp_events_total{outcome=ok|late|duplicate|dlq}, fp_ingest_latency_seconds (vehicle timestamp → processed), fp_alerts_total{rule}, fp_batch_seconds, fp_api_latency_seconds{route}. docker compose --profile observability starts Prometheus and Grafana with a provisioned “FleetPulse pipeline” dashboard (events/s by outcome, ingest latency p50/p95/p99, alerts by rule, batch time, API p95 by route, Kafka ingest rate). Consumer lag comes from rpk group describe / the Kafka exporter. Logs are structured stdout collected by the platform (Loki/ELK). The dashboard itself shows live events/s and measured alert latency.



Troubleshooting a latency spike: (1) Grafana: is fp_api_latency p95 up for one route or all? (2) If all, check pod CPU and the Postgres connection pool; if one route, run EXPLAIN ANALYZE on its query (as in 5.3). (3) For alert latency, compare fp_ingest_latency with consumer lag: rising lag means processors are saturated, so scale replicas up to the partition count; flat lag but slow batches (fp_batch_seconds) points to Postgres writes. (4) Correlate with simulator burst logs. OpenTelemetry tracing across services is the next step.

11. AI / ML Component (if used)

- **Purpose:** The decision the model or agent supports, and why rules alone were not enough
- **Data & Features:** Training or retrieval data, how it was generated or sourced, leakage checks
- **Model / Agent Design:** Model choice, vector store, tools the agent can call, how access is limited
- **Evaluation:** Metrics against a simple baseline
- **Guardrails & Cost:** Failure handling, hallucination checks, human-in-the-loop, cost and latency per request

Purpose: rank every vehicle by its probability of breaking down in the next 7 days. Rules alone catch faults that have already happened; the model combines weak, noisy signals (slowly rising coolant, sagging 12 V voltage, an uptick in minor codes, overdue service) into an early warning.

Data and features: 100,000 simulated vehicle-weeks from the same wear model as the live simulator; features: service-overdue ratio, age, odometer, max coolant (7 d), min 12 V (7 d), fault-code events (7 d), harsh brakes (7 d), EV / hybrid flags. Live features come from the vehicle_daily rollup, which the batch job maintains incrementally (only rows since the last watermark are read, each counted exactly

once). Leakage check: the label is drawn from a hazard on the hidden wear, which is never a feature; only noisy symptoms are.

Model: scikit-learn HistGradientBoostingClassifier (200 iterations); per-vehicle explanations from importance-weighted deviation from the fleet median. Vector store: pgvector HNSW over a fault/repair knowledge base, embedded with a dependency-free hashed n-gram embedding so it runs offline.

Metric (held-out 25%)	Model	Baseline: km since service	Lift
ROC-AUC	0.878	0.688	+0.19
PR-AUC	0.190	0.054	3.5x
Precision @ top 2%	27.8%	6.6%	4.2x

Agent: six tools (top_risk, vehicle_summary, open_alerts, search_knowledge, estimate_savings, propose_work_order), all filtered by the caller's tenant. Planner: Claude tool-use if ANTHROPIC_API_KEY is set, else a deterministic intent router. Guardrails: injection screen, max 6 tool calls, VIN grounding check, propose-only actions requiring human approval, every tool call audited as agent:<user>. Cost/latency: deterministic mode ~20–40 ms and US\$0; LLM mode is a few model calls per question.

12. Architecture Decisions, Risks & Future Enhancements

- **Architecture Decision Records:** 3–5 records: Context → Options Considered → Decision → Consequences (cover databases, messaging and one CAP trade-off)
- **Risks & Technical Debt:** Known limitations and shortcuts taken during the hackathon
- **Future Enhancements:** Next three steps from prototype to pilot

ADRs (docs/adr): 0001 Kafka as telemetry backbone (replay, per-VIN ordering, DLQ); 0002 polyglot persistence (Postgres 3NF + day-partitioned telemetry + pgvector, Redis for hot state); 0003 CAP/PACELC per data class and effectively-once delivery; 0004 single Python image with batch hot path and a deterministic-first copilot.

Risks and shortcuts: 100K events/s not proven end to end: on one laptop the pipeline sustains ~24K events/s with zero loss and Kafka absorbs 69K/s, with a single Postgres primary as the ceiling; 100K/s needs the telemetry store beyond one Postgres primary (ADR-0002). mTLS, OIDC and Vault are designed, not running. The model is validated on simulated data only. Terraform is unapplied. Bloom state is per process and lost on restart (DB constraints keep correctness).

Next three steps to pilot: (1) move raw telemetry to ClickHouse or Scylla with Parquet cold storage and run a 100K events/s k6/Kafka load test on EKS; (2) Avro + schema registry and mTLS device identity; (3) retrain on a real pilot fleet's maintenance history and add survival-analysis time-to-failure.

13. Demo Video (5 Minutes Maximum)

Videos longer than 5:00 are cut at 5:00. Show the working product, not only slides. 1080p, clear audio, captions recommended; unlisted YouTube, Vimeo or Drive link.

Time	Segment	What to Show
0:00 – 0:35	Problem	US\$2,400 per breakdown, ~2,800 a week at 100K vehicles; FleetPulse predicts and flags faults in under a second
0:35 – 1:15	Architecture	Simulator → Kafka → processors → Postgres / Redis / pgvector; batch rollup + ML scoring
1:15 – 2:00	Live dashboard	Manager login; alert latency (~0.2 s) and data freshness (~0.15 s) tiles; map; live critical alerts
2:00 – 3:05	Prediction + copilot	Top-risk vehicle with reasons; 0.88 vs 0.69 ROC-AUC; savings estimate; prompt injection refused; audit
3:05 – 3:55	Crash and recover (live)	Kill all processors + restart the broker; recovers in ~30 s, zero loss, no duplicates
3:55 – 4:25	Observability	Grafana pipeline dashboard showing the outage and recovery
4:25 – 5:50	Evidence	Green CI (unit 98%, integration, contract, BDD, ZAP, Trivy, chaos, K8s on kind); 25K/s zero-loss load test; API p95 110 ms @50; 45-min soak
5:50 – 6:30	Close	Impact, next steps (ClickHouse for 100K/s, SSO, device certificates)

- Video Link: https://drive.google.com/file/d/1R0No_tNViiExbQ-IPfvOw9y1ivHt7P9Z/view?usp=sharing (script: docs/demo-script.md)

14. Repository Checklist

- **README:** Problem, architecture image, quick start, environment variables, test commands, known issues
- **One-Command Run:** docker compose up (or make run) starts the full stack with the data simulator
- **Structure:** Clear service folders, plus /docs, /infra and /tests
- **CI Pipeline:** Build, lint, all test suites and security scans on every push
- **Hygiene:** No secrets committed, .env.example provided, commits from all members
- **Final Tag:** v1.0-submission

15. Conclusion

- Summary of key learnings

- Strengths of the solution
- Challenges tackled and solved

Key learnings: idempotent sinks make at-least-once delivery safe and are far simpler than distributed transactions; partitioning by time turns retention into a metadata operation, but only if the rollups are incremental (our first materialised-view version re-scanned all history every minute and grew ~5 GB/hour until we replaced it); most API latency wins came from indexes that match the query's filter and order.

Strengths: a working end-to-end system on 100,000 vehicles with measured numbers, an ML model that clearly beats the status-quo rule, and security and compliance built into the data model rather than added on.

Challenges solved: duplicate and out-of-order events, two incompatible OEM formats, per-tenant isolation from the database to the WebSocket, and recovery after killing processors and the broker.

16. Declarations

- **Open-Source Components:** List with licences (an SBOM file in the repository is acceptable)
- **AI Tools Used:** Coding assistants or LLM APIs, and the parts of the work they were used for
- **Data:** Confirm all data is synthetic or public and contains no real personal data

Open-source components: FastAPI (MIT), Uvicorn (BSD), Pydantic (MIT), psycpg 3 (LGPL), redis-py (MIT), confluent-kafka / librdkafka (Apache-2.0 / BSD), scikit-learn (BSD), NumPy (BSD), joblib (BSD), PyJWT (MIT), prometheus-client (Apache-2.0), anthropic SDK (MIT), Leaflet (BSD-2), pytest-bdd (MIT), Esri World Dark Gray Canvas map tiles (Esri terms, attributed in the UI), PostgreSQL (PostgreSQL licence), pgvector (PostgreSQL licence), Redis 7 (BSD), Redpanda (BSL, dev only), Prometheus and Grafana (Apache-2.0 / AGPL).

AI tools used: Claude (Anthropic) via Claude Code was used to plan, write and test code, generate diagrams and draft this document; the team reviewed and owns the result. The running copilot can optionally call the Anthropic Messages API.


Data: all vehicles, VINs, drivers, positions and fault histories are synthetic and generated by our simulator; no real personal or vehicle-owner data is used.

17. Appendix (if any)

- Additional diagrams or flowcharts
- Full test logs or large datasets used
- References or external sources
- Evidence files in the repository: docs/evidence/api_load.txt, processor_bench.txt, alert_latency.txt, soak.txt, grafana.png, explain_analyze.txt, chaos.txt, dashboard.png,

vehicle_drawer.png; fleetpulse/ml/metrics.json; docs/threat-model.md; docs/adr/*.md;
docs/demo-script.md.

Relational core (3NF) - arrows point from child (FK) to parent



- Telemetry (partitioned by hour, unique (vin, seq, ts)), the vehicle_daily incremental rollout and fault_knowledge (pgvector) sit outside the 3NF core. Deliberate denormalisation: alert.tenant_id and risk_score.tenant_id avoid a 3-table join on the hottest tenant-filtered queries.

- References: Connected Vehicle Intelligence Hackathon problem statement (Talenciaglobal, 2026); SAE J2012 (DTC definitions); ISO 3779 / NHTSA 49 CFR 565 (VIN check digit); Kafka documentation (idempotent producer, consumer groups); PostgreSQL 16 declarative partitioning; pgvector HNSW.